

Recommendation Systems - Rating Predictions with MovieLens Data

Danilo Ferreira de Oliveira

04/22/2021

Contents

1	Introduction	2
2	Analysis	2
2.1	Getting the data	2
2.2	Exploratory data analysis and visualization	3
2.2.1	Exploring the variables	3
2.2.2	Ratings	4
2.2.3	Movie and user IDs	5
2.2.4	Release date	6
2.2.5	Genres and genre combinations	8
2.2.6	Rating date	11
2.3	Feature engineering	13
2.4	Creating training and test sets	15
2.5	Prediction models	15
2.5.1	Overall Average Rating	16
2.5.2	Movie effect model	16
2.5.3	Movie + User effects model	17
2.5.4	Movie + User + Genre effects model	18
2.5.5	Regularized Movie + User effects model	20
2.5.6	Matrix Factorization with <i>recosystem</i> package	22
3	Results	23
3.1	Training full edx dataset with Matrix Factorization	23
3.2	Validation set's predicted ratings and RMSE	23
4	Conclusion	24

1 Introduction

For this project, we will be creating a movie recommendation system using the MovieLens dataset. You can find the [entire latest MovieLens dataset here](#). We will be creating your own recommendation system using all the tools shown throughout the courses in the *HarvardX: Data Science Professional Certificate* series. We will use the [10M version of the MovieLens dataset](#) to make the computation a little easier.

We will train machine learning algorithm using the inputs in one subset to predict movie ratings in the validation set.

Recommendation systems are more complicated machine learning challenges because each outcome has a different set of predictors. For example, different users rate a different number of movies and rate different movies. By evaluating the movies specific users like to watch, we can predict ratings they would give for movies they haven't seen, and then recommend them those with relatively high ratings.

We will firstly follow the approach made by Professor Rafael Irizarry in [Introduction to Data Science](#), and continue it with the addition of genre effects, and a model that uses matrix factorization.

To compare different models or to see how well we are doing compared to a baseline, we will use root mean squared error (RMSE) as our loss function. We can interpret RMSE similar to standard deviation.

We will use a script provided by HarvardX to generate our datasets. After this, we analyze the obtained training dataframe and develop algorithms to predict ratings. For a final test of our final algorithm, we predict movie ratings in the validation set (the final hold-out test set) as if they were unknown. RMSE will be used to evaluate how close your predictions are to the true values in the final hold-out test set.

2 Analysis

2.1 Getting the data

Before we can start gathering our data, we load the necessary R libraries for the development of the whole report.

```
library(caret)
library(tidyverse)
library(ggplot2)
library(lubridate)
library(stringr)
library(recosystem)
library(data.table)
library(kableExtra)
library(tibble)
```

Next, we download the *MovieLens* data and run code provided by *edX* to generate the datasets for training and evaluating the final model. The `edx` set will be used all along the report to train and choose the best model for rating predictions, while the `validation` set will only be used at the end of to evaluate the quality of the final model with the root mean square error.

```
dl <- tempfile()
download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)

ratings <- fread(text = gsub("::", "\t", readLines(unzip(dl, "ml-10M100K/ratings.dat"))),
                 col.names = c("userId", "movieId", "rating", "timestamp"))
```

```

movies <- str_split_fixed(readLines(unzip(dl, "ml-10M100K/movies.dat")), "\\: ", 3)
colnames(movies) <- c("movieId", "title", "genres")

movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(movieId),
                                           title = as.character(title),
                                           genres = as.character(genres))

movielens <- left_join(ratings, movies, by = "movieId")

```

The validation set will be 10% of the *MovieLens* data obtained. We need to make sure all user and movie IDs in validation are also in `edx`, so we apply `semi_join()` transformations. Then, we add rows removed from the validation set back into the training set.

```

set.seed(1, sample.kind="Rounding")
test_index <- createDataPartition(y = movielens$rating, times = 1, p = 0.1, list = FALSE)
edx <- movielens[-test_index,]
temp <- movielens[test_index,]

validation <- temp %>%
  semi_join(edx, by = "movieId") %>%
  semi_join(edx, by = "userId")

removed <- anti_join(temp, validation)
edx <- rbind(edx, removed)

rm(dl, ratings, movies, test_index, temp, movielens, removed)

```

We can now start analyzing the `edx` dataframe.

2.2 Exploratory data analysis and visualization

2.2.1 Exploring the variables

To begin, let's take a look at our dataset columns and variables.

```
edx %>% as_tibble()
```

```

>> # A tibble: 9,000,055 x 6
>>   userId movieId rating timestamp title          genres
>>   <int>   <dbl>   <dbl>     <int> <chr>         <chr>
>> 1     1     122     5 838985046 Boomerang (1992) Comedy|Romance
>> 2     1     185     5 838983525 Net, The (1995) Action|Crime|Thriller
>> 3     1     292     5 838983421 Outbreak (1995) Action|Drama|Sci-Fi|T-
>> 4     1     316     5 838983392 Stargate (1994) Action|Adventure|Sci--
>> 5     1     329     5 838983392 Star Trek: Generation~ Action|Adventure|Dram~
>> 6     1     355     5 838984474 Flintstones, The (199~ Children|Comedy|Fanta~
>> 7     1     356     5 838983653 Forrest Gump (1994) Comedy|Drama|Romance|~
>> 8     1     362     5 838984885 Jungle Book, The (199~ Adventure|Children|Ro~
>> 9     1     364     5 838983707 Lion King, The (1994) Adventure|Animation|C~
>> 10    1     370     5 838984596 Naked Gun 33 1/3: The~ Action|Comedy
>> # ... with 9,000,045 more rows

```

It is possible to see that columns `title` and `genres` are of character class (<chr>), while `movieId` and `rating` are numeric (<dbl>), `userId` and `timestamp` are integers (<int>).

We should analyze the quantity of unique values for each important feature column.

```
edx %>% summarize(n_userIds = n_distinct(userId), n_movieIds = n_distinct(movieId),
                  n_titles = n_distinct(title), n_genres = n_distinct(genres)) %>%
  kbl(booktabs = TRUE,
      caption = "Distinct values for userId, movieId, titles and genres") %>%
  kable_styling(latex_options = c("striped", "HOLD_position"))
```

Table 1: Distinct values for userId, movieId, titles and genres

n_userIds	n_movieIds	n_titles	n_genres
69878	10677	10676	797

Note that there are more `movieId` unique variables than there are `title` ones. We need to investigate why.

```
edx %>% group_by(title) %>%
  summarize(n_movieIds = n_distinct(movieId)) %>% filter(n_movieIds > 1) %>%
  kbl(booktabs = TRUE, caption = "Movies with more than one ID") %>%
  kable_styling(latex_options = c("striped", "HOLD_position"))
```

Table 2: Movies with more than one ID

title	n_movieIds
War of the Worlds (2005)	2

So we identify that *War of the Worlds (2005)* has two distinct `movieId` numbers. Let's see how many reviews each of them have.

Table 3: Different values of movieId and genres for *War of the Worlds (2005)*

movieId	title	genres	n
34048	War of the Worlds (2005)	Action Adventure Sci-Fi Thriller	2460
64997	War of the Worlds (2005)	Action	28

The second `movieId` received only 28 reviews. This is a very low value, and since there is a possibility both IDs are in the validation set, we won't meddle with it.

2.2.2 Ratings

We can now visualize the data. Let's see how the ratings distribution is configured.

From the plot above, we can affirm that full star ratings are more common than those with half stars, since 4, 3 and 5 are the most frequent ones.




Figure 1: Distribution of ratings

2.2.3 Movie and user IDs

Let's evaluate the frequency distribution of reviews that each user gives and reviews each movie receives.

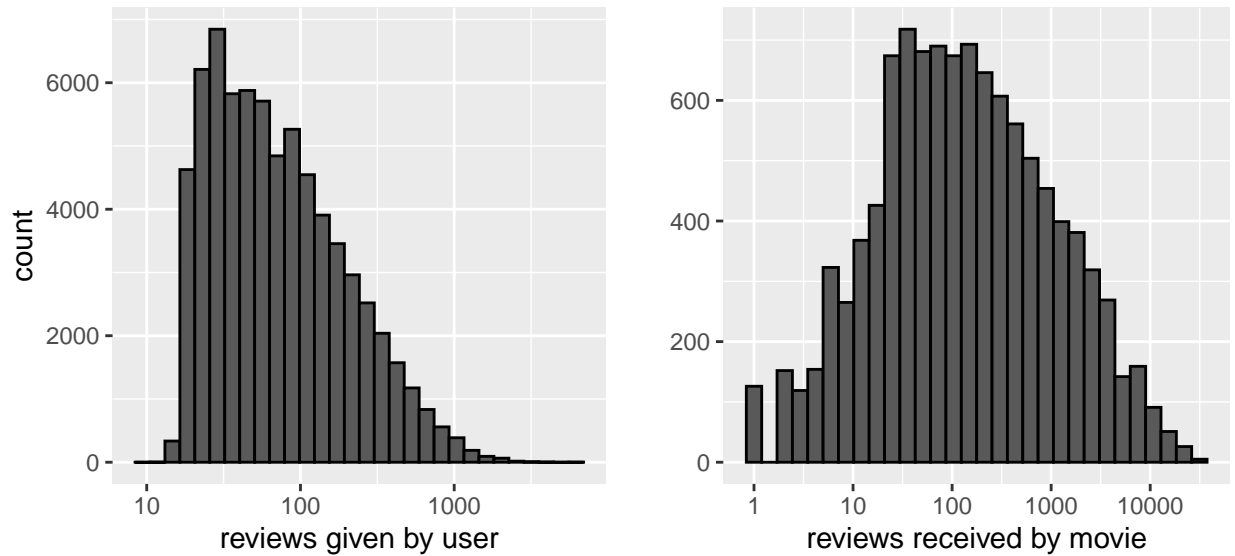


Figure 2: Frequency of reviews given by a specific user (left) and of reviews received by a specific movie (right)

The x-axis from the plots above are both on a logarithmic scale, and the one on the left is **rightly skewed**. If we analyze them, we can say that the majority of the users give from 20 to 200 reviews, while each movie frequently receive between 20 to a 1000 reviews, which is a pretty broad range.

Plotting histograms for average ratings, we get the figures below. From the top plot, it is possible to see that it is rare for a user to average rating movies below 2, since there are few users to the left of that number (only 185 of all 69878 in the dataset, to be more exact).

From all 10,676 movies, we can see that only 61 of them average below 1.5, and 57 films average over 4.25

rating.




Figure 3: Frequency of user average ratings (top) and movie average ratings (bottom)

Let's now explore the top 10 most reviewed movies in the dataset and how many times they were reviewed.

As we can see, movies that receive the most reviews are normally blockbusters. Checking for their average ratings, we also notice that they are really high, the lowest of them being 3.66. We can infer that, the more a movie is reviewed, the better its rating is. It makes sense, being the better the movie, the higher its ratings and the more people watch it, subsequently more people can rate them.

the title of the movies come along with the release date, inside parenthesis.

2.2.4 Release date

Since there isn't a column for each movie's release date, we will check it by using the `str_extract` function from the `stringr` package.

```
release_date <- edx$title %>% str_extract('\\([0-9]{4}\\)') %>%
  str_extract('[0-9]{4}') %>% as.integer()
summary(release_date)
```

```
>>   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
>>  1915   1987   1994   1990   1998   2008
```

We first apply this function extracting the date along with the parenthesis, and then the number from inside them. We do it like this because by extracting only the number, movies like *2001: A Space Odyssey* (1968) and *2001 Maniacs* (2005) show up when searching for year 2001, for example. To go even further, it also extracts the numbers 1000 and 9000 for *House of 1000 Corpses* (2003) and *Detroit 9000* (1973), respectively.

Below we can see that the earliest release date in our dataset (1915) only has one movie:

Table 4: Movie with the earliest release date in the dataset

movieId	title	genres	n
7065	Birth of a Nation, The (1915)	Drama War	180

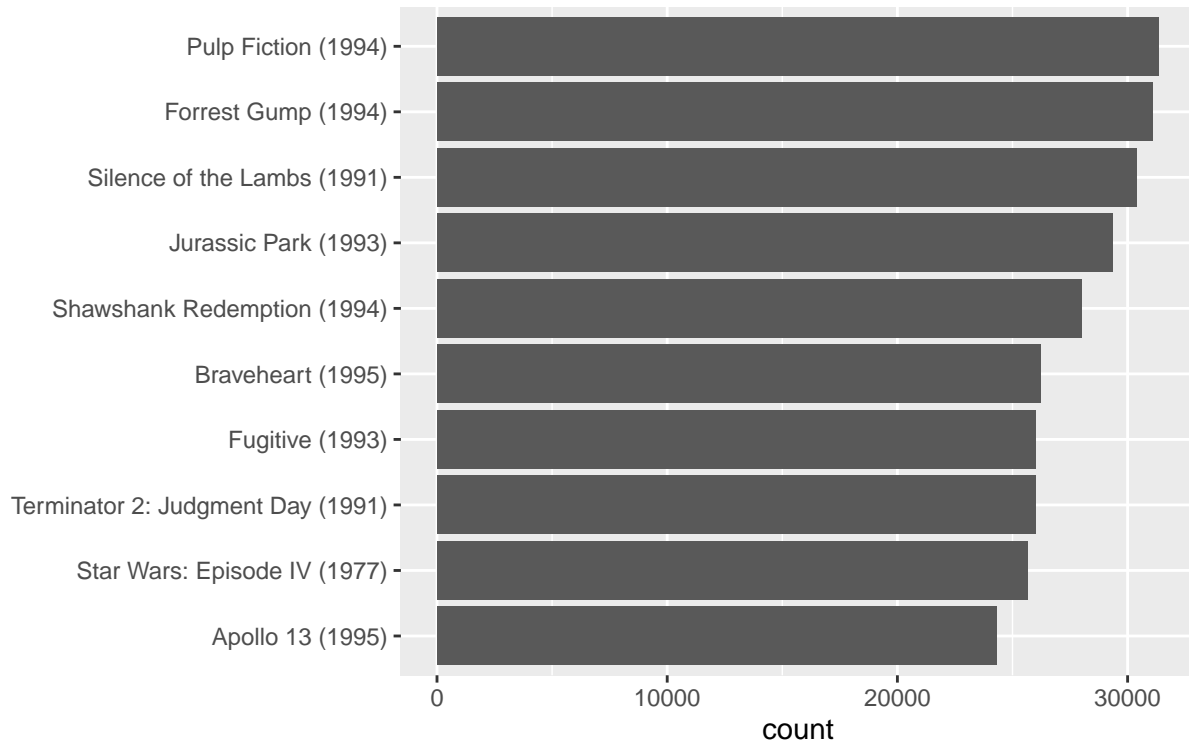


Figure 4: Number of reviews for the top 10 most reviewed movies




Figure 5: Average rating for the top 10 most reviewed movies

We now compute the number of ratings for each movie and then plot it against the year the movie came out, using a boxplot for each year. The logarithmic transformation is applied on the y-axis (number of ratings) when creating the plot.




Figure 6: Number of reviews the movies received versus the year they were released

We see that, on average, movies that came out in the mid 90's get more ratings, reaching values over 30,000 reviews. We also see that with newer movies, starting near 1998, the number of ratings decreases with year: the more recent a movie is, the less time users have had to rate it.

Now we compute the average of ratings for each movie and plot it against the release year with boxplots, similarly to before.

It is possible to see that most movies released until 1973 normally range between 3.0 and 4.0 ratings. From then on, the range gets bigger and values get lower. The median value goes from averaging over 3.5 to near 3.25. There are even some movies from around the 2000's that average 0.5 ratings, but they are outliers and probably only received very few ratings.

2.2.5 Genres and genre combinations

The **genres** variable actually represents combinations of a series of unique genres. We will count the number of reviews for all different combinations that appear in our dataset and arrange them in descending order. The table below shows, as an example, the top 7 most reviewed genre combinations.




Figure 7: Average rating the movies received versus the year they were released

Table 5: Top genre combinations rated

rank	genres	count
1	Drama	733296
2	Comedy	700889
3	Comedy Romance	365468
4	Comedy Drama	323637
5	Comedy Drama Romance	261425
6	Drama Romance	259355
7	Action Adventure Sci-Fi	219938

Let's check the genre effect on ratings. We will filter for genre combinations that were reviewed over a thousand times, put them in ascending order of average rating and show 15 of them, equally distant inside the arranged dataframe. The indexes to select the genres in the arranged set are represented by the variable `gcomb_15`.

```
gcomb_15 <- c(1,seq(31,415,32),444)
edx %>% group_by(genres) %>%
  summarize(n = n(), avg = mean(rating), se = sd(rating)/sqrt(n())) %>%
  filter(n > 1000) %>% arrange(desc(avg)) %>% .[gcomb_15,] %>%
  mutate(genres = reorder(genres, avg)) %>%
  ggplot(aes(x = genres, y = avg,
             ymin = avg - 2*se, ymax = avg + 2*se)) + geom_point() +
  geom_errorbar() + theme(axis.text.x=element_text(angle = 30, hjust = 1)) +
  labs(x='', y = 'average rating')
```




Figure 8: Distribution of average ratings for a few genre combinations possible

From the plot above, it is clear the effect genre combinations have on the rating, since they range from approximately 2.0 to over 4.5. Let's now see the effect of each genre separately.

There are 20 different genres in this dataset, and they are the following:

Table 6: All unique genres present in the dataset

Comedy	Drama	War	Film-Noir
Romance	Sci-Fi	Animation	Horror
Action	Adventure	Musical	Documentary
Crime	Children	Western	IMAX
Thriller	Fantasy	Mystery	(no genres listed)

We can count how many reviews each genre have had, by applying the following function:

```
all_genres_count <- sapply(all_genres, function(g) {
  sum(str_detect(edx$genres, g))
}) %>% data.frame(review_count=.) %>%
  arrange(desc(.))
```

We can see the results on the table below. The genres range from near 4 million to only 7 reviews.

Table 7: Number of reviews per genre

all_genres_count[1:10]		all_genres_count[11:20]	
Genre	Count	Genre	Count
Drama	3910127	Horror	691485
Comedy	3540930	Mystery	568332
Action	2560545	War	511147
Thriller	2325899	Animation	467168
Adventure	1908892	Musical	433080
Romance	1712100	Western	189394
Sci-Fi	1341183	Film-Noir	118541
Crime	1327715	Documentary	93066
Fantasy	925637	IMAX	8181
Children	737994	(no genres listed)	7

We need to investigate the (no genres listed) genre, which actually represents no genre at all. We can see the movies that belong to it and their respective averages and standard errors:

```
edx[which(edx$genres=='(no genres listed)'),] %>%
  summarize(movieId=movieId[1], userId=userId[1], title=title[1], genres=genres[1],
            n = n(), avg = mean(rating), se = sd(rating)/sqrt(n()))
```

```
>>  movieId  userId          title          genres n      avg      se
>> 1    8606    7701 Pull My Daisy (1958) (no genres listed) 7 3.642857 0.4185332
```

There is only one movie in the dataset with no genres listed, and it has only 7 reviews, giving its average rating a high standard error (0.4185). For this reason, we won't be plotting this genre's average rating along with the others.

The plot below shows the average ratings for each unique genre, as they range from under 3.3 to over 4.0.

We now observe the top genres reviewed in this dataset, and how many reviews each one of them have. We accounted for the genres that were reviewed over a million times, resulting in 8 distinct possibilities. The reviews count for all genre has already been shown, but we will plot the count for the top 8 genres to better visualize it, and also compare their shape with the movie count plot for the top 8.

We now plot how many different movies were reviewed for each of the top 8 genres.

The general shape and proportion of the bars pretty much look the same for both reviews and movie counts.

2.2.6 Rating date

Since `timestamp` doesn't really have any value for us as it is, we create from it another column that gives us the date and time a movie was rated.

```
edx <- mutate(edx, date = as_datetime(timestamp))
edx %>% as_tibble()
```

```
>> # A tibble: 9,000,055 x 7
>>   userId movieId rating timestamp title          genres          date
>>   <int>   <dbl>   <dbl>     <int> <chr>          <chr>          <dtm>
```




Figure 9: Average rating for each genre

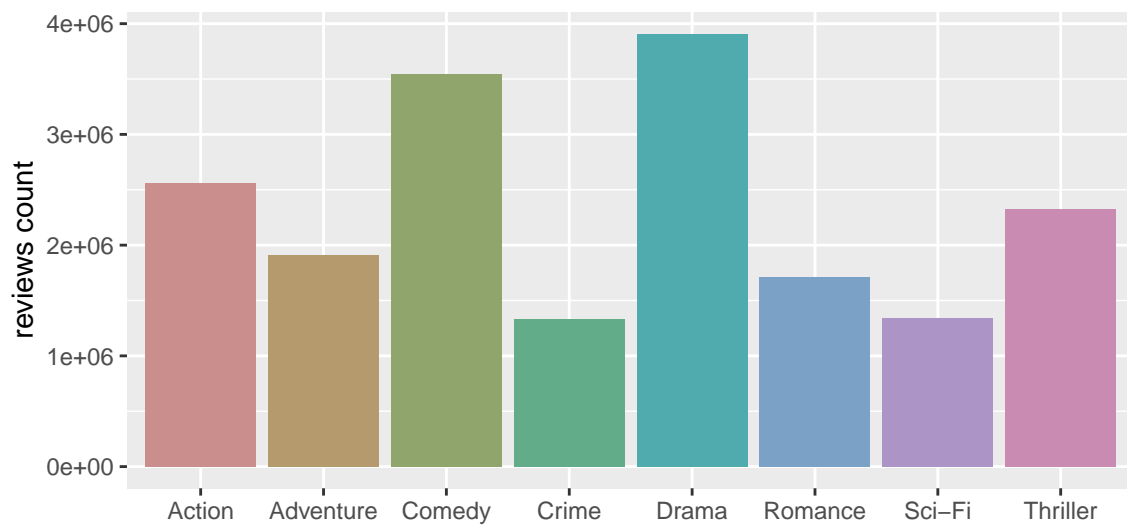


Figure 10: Number of reviews for each of the top 8 most reviewed genres




Figure 11: Number of movies that belongs to the top 8 most reviewed genres

```
>> 1      1      122      5 838985046 Boomerang (~ Comedy|Roma~ 1996-08-02 11:24:06
>> 2      1      185      5 838983525 Net, The (1~ Action|Crim~ 1996-08-02 10:58:45
>> 3      1      292      5 838983421 Outbreak (1~ Action|Dram~ 1996-08-02 10:57:01
>> 4      1      316      5 838983392 Stargate (1~ Action|Adve~ 1996-08-02 10:56:32
>> 5      1      329      5 838983392 Star Trek: ~ Action|Adve~ 1996-08-02 10:56:32
>> 6      1      355      5 838984474 Flintstones~ Children|Co~ 1996-08-02 11:14:34
>> 7      1      356      5 838983653 Forrest Gum~ Comedy|Dram~ 1996-08-02 11:00:53
>> 8      1      362      5 838984885 Jungle Book~ Adventure|C~ 1996-08-02 11:21:25
>> 9      1      364      5 838983707 Lion King, ~ Adventure|A~ 1996-08-02 11:01:47
>> 10     1      370      5 838984596 Naked Gun 3~ Action|Come~ 1996-08-02 11:16:36
>> # ... with 9,000,045 more rows
```

The new column `date` is of a Date-Time class (`<dtm>`) called `POSIXct`, and it clearly makes it better to understand the time of rating.

We can now plot the time effect on the dataset, showing the influence the week a movie was rated on the average value. Along with ratings, we plot a LOESS smooth line to better show the relationship between variables.

```
edx %>% mutate(date = round_date(date, unit = "week")) %>%
  group_by(date) %>%
  summarize(rating = mean(rating)) %>%
  ggplot(aes(date, rating)) +
  geom_point() + labs(y='average rating') +
  geom_smooth()
```

The review date definitely has some effect on ratings, but it isn't significant.

2.3 Feature engineering

Down the line, we will need all genres available in the dataset as binary values, so it is necessary to create columns for all 20 of them. For each row, we apply value 1 to their corresponding genres in the character column, and value 0 to the rest. here we see the first 6 rows, with only the `movieId`, title and genres of the movies, so we can compare the genres further in the analysis.




Figure 12: The review date effect on the average rating

```
edx[,c(2,5,6)] %>% as_tibble(.rows = 6)
```

```
>> # A tibble: 6 x 3
>>   movieId title                               genres
>>   <dbl> <chr>                                <chr>
>> 1    122 Boomerang (1992)          Comedy|Romance
>> 2    185 Net, The (1995)          Action|Crime|Thriller
>> 3    292 Outbreak (1995)          Action|Drama|Sci-Fi|Thriller
>> 4    316 Stargate (1994)          Action|Adventure|Sci-Fi
>> 5    329 Star Trek: Generations (1994) Action|Adventure|Drama|Sci-Fi
>> 6    355 Flintstones, The (1994) Children|Comedy|Fantasy
```

We apply the transformation by doing the following `for()` command. We will also permanently remove columns such as `timestamp`, `title`, `genres` and `date`, since we won't be needing them to train the prediction models.

```
for (var in all_genres) {
  edx <- edx %>% mutate(genre=ifelse(str_detect(genres,as.character(var)),1,0))
  colnames(edx)[ncol(edx)] <- as.character(var)
}
rm(var)
edx <- edx[,-(4:7),drop=FALSE]
edx[,4:23] %>% as_tibble(.rows=6)
```

```
>> # A tibble: 6 x 20
>>   Comedy Romance Action Crime Thriller Drama `Sci-Fi` Adventure Children Fantasy
```

```

>>      <dbl>      <dbl>      <dbl> <dbl>      <dbl> <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
>> 1          1          1          0    0          0    0          0          0          0          0
>> 2          0          0          1    1          1    0          0          0          0          0
>> 3          0          0          1    0          1    1          1          0          0          0
>> 4          0          0          1    0          0    0          1          1          0          0
>> 5          0          0          1    0          0    1          1          1          0          0
>> 6          1          0          0    0          0    0          0          0          1          1
>> # ... with 10 more variables: War <dbl>, Animation <dbl>, Musical <dbl>,
>> #   Western <dbl>, Mystery <dbl>, Film-Noir <dbl>, Horror <dbl>,
>> #   Documentary <dbl>, IMAX <dbl>, (no genres listed) <dbl>

```

Comparing each row in this tibble above with the ones right before, it is possible to see that the genres in the first one are represented as 1s in the other, while genres that don't appear before are kept as zero values. The transformation occurred accurately.

2.4 Creating training and test sets

Before training any model, we need to create a train and a test set. We first set the seed to guarantee the reproducibility of the model, and we divide the `edx` dataframe into 80-20% partitions.

```

set.seed(123, sample.kind = "Rounding")
index <- createDataPartition(edx$rating, times=1, p=0.2, list= FALSE)
test <- edx[index,]
train <- edx[-index,]

```

To make sure all user and movie IDs in `test` are also in `train`, we apply the `semi_join()` function, just as we did in the beginning of the report for the first two datasets.

```

test <- test %>%
  semi_join(train, by = "movieId") %>%
  semi_join(train, by = "userId")

```

2.5 Prediction models

Now that we have our training and test sets, we can start evaluating prediction models. We will begin by following the approach in Prof. Irizarry's book, starting with just a Overall Average Rating model, going through the Movie effect, Movie + User effect and Regularized Movie + User effects models. We will also apply a method that includes Genre effect and one that performs Matrix Factorization. To evaluate their quality, we will use the RMSE function.

If N is the number of user-movie combinations, $y_{u,i}$ is the rating for movie i by user u , and $\hat{y}_{u,i}$ is our prediction, then RMSE is defined as follows:

$$\text{RMSE} = \sqrt{\frac{1}{N} \sum_{u,i} (\hat{y}_{u,i} - y_{u,i})^2}$$

which is represented by the function:

```

RMSE <- function(true_ratings, predicted_ratings){
  sqrt(mean((true_ratings - predicted_ratings)^2))
}

```

2.5.1 Overall Average Rating

We start with a model that assumes the same rating for all movies and all users, with all the differences explained by random variation: If μ represents the true rating for all movies and users and ε represents independent errors sampled from the same distribution centered at zero, then:

$$Y_{u,i} = \mu + \varepsilon_{u,i}$$

In this case, the least squares estimate of μ — the estimate that minimizes the root mean squared error — is the average rating of all movies across all users.

```
mu <- mean(train$rating)
```

In this first model, we will predict all ratings as the average rating of the training set and compare it to the actual ratings in the test set (`RMSE(test$rating, mu)`).

Table 8: RMSE result for the Average Method

Method	RMSE
Just the Average	1.060013

We can observe that this value is actually the standard deviation of this distribution. Its value is too high, thus not acceptable.

$$\text{RMSE}_\mu = \sqrt{\frac{1}{N} \sum_{u,i} (\mu - y_{u,i})^2} = \sigma$$

2.5.2 Movie effect model

We can improve our model by adding a term, b_i , that represents the effect of each movie's average rating:

$$Y_{u,i} = \mu + b_i + \varepsilon_{u,i}$$

Here, b_i is approximately the average of $Y_{u,i}$ minus the overall mean for each movie i . We can again use least squares to estimate the b_i in the following way:

```
fit <- lm(rating ~ as.factor(movieId), data=edx)
```

Note that because there are thousands of b 's, the `lm()` function will be very slow or cause R to crash, so it is not recommendable using linear regression to calculate these effects. So we will calculate all b_i 's, and subsequently the predicted ratings, in the way shown below.

```
movie_avgs <- train %>% group_by(movieId) %>% summarize(b_i = mean(rating-mu))

predicted_ratings <- mu + test %>% left_join(movie_avgs, by='movieId') %>%
  pull(b_i)

summary(predicted_ratings)
```



```
>>   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
>>  0.500   3.217   3.586   3.512   3.878   4.750
```

From the summary above, we can see that the predictions obtained range from 0.5 and 4.75, inside the real range (0.5 to 5). Calculating `RMSE(test$rating, predicted_ratings)`, we get the results shown in the next table.

Table 9: RMSE results with the Movie effect model

Method	RMSE
Just the Average	1.0600131
Movie effect model	0.9436204

A 12% drop is a big improvement, but we can still do better.

2.5.3 Movie + User effects model

We further improve our model by adding b_u , the user-specific effect:

$$Y_{u,i} = \mu + b_i + b_u + \varepsilon_{u,i}$$

To fit this model, we could again calculate the least squares estimate by using `lm` like this:

```
lm(rating ~ as.factor(movieId) + as.factor(userId))
```

But for the same reason explained before, we will follow a simpler model approach, with b_u being calculated with the difference between $Y_{u,i}$ and both overall mean rating and movie-specific coefficient. It is demonstrated in the code chunk below:

```
user_avgs <- train %>% left_join(movie_avgs, by='movieId') %>%
  group_by(userId) %>% summarize(b_u = mean(rating - mu - b_i))

predicted_ratings <- test %>%
  left_join(movie_avgs, by='movieId') %>% left_join(user_avgs, by='userId') %>%
  mutate(pred = mu + b_i + b_u) %>% pull(pred)

summary(predicted_ratings)
```

```
>>   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
>> -0.7161  3.1353  3.5651  3.5127  3.9450  6.0547
```

Predicted ratings now have a spread from -0.7161 to 6.0547, which is bigger than the real range (0.5 to 5). There are 208 values under 0.5 and 4149 over 5. But what we really need to evaluate is the RMSE value by applying the RMSE function.

Table 10: RMSE results with the Movie and User effects model

Method	RMSE
Just the Average	1.0600131
Movie effect model	0.9436204
Movie + User effects model	0.8662966

Another big improvement in the RMSE value with a 8% drop. Let's proceed with other models.

2.5.4 Movie + User + Genre effects model

Now we try an approach proposed by Prof. Irizarry, which was presented in his book but not built and developed in it, that adds the genre effect to the model. Its formula is presented as shown below:

$$Y_{u,i} = \mu + b_i + b_u + \sum_{k=1}^K x_{u,i}^k \beta_k + \varepsilon_{u,i}$$

with $x_{u,i}^k = 1$ if $g_{u,i}$ is genre k . We already determined $x_{u,i}$ values earlier by creating binary columns for each of the 20 genres available in the training and test data. Now we determine β_k values with the same approach utilized in the last two models, where each β_k is the mean of the subtraction of overall average rating, movie-specific coefficient b_i and user-specific coefficient b_u from ratings $Y_{u,i}$. This way, if a movie belongs to both Comedy and Drama, one β will throw the average rating up (high ratings for Drama) and the other down (low average for Comedies).

To make all the process faster down the line, we will first add b_i and b_u columns to the training and test sets, and then we calculate all β_k 's.

```
train <- train %>% left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId')
test <- test %>% left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId')
rm(user_avgs, movie_avgs)

beta_k <- vector()
for (i in seq_along(all_genres)) {
  b_value <- train %>%
    group_by(!!sym(all_genres[[i]])) %>%
    summarize(beta_k=mean(rating-mu-b_i-b_u)) %>%
    filter((.[1])==1) %>% .[[2]]
  beta_k <- append(beta_k, b_value)
}
rm(i, b_value)

df_beta <- data.frame(beta_k)
rownames(df_beta) <- all_genres
```

First, we group the training set by each genre column by unquoting their names with argument `!!` and function `sym`, which forms two different groups (genre = 0 and genre = 1). We then calculate each specific genre's beta-values, and filter for column values equal to 1 (the genre is present). We append it to the `beta_k` vector and, lastly, transform it into a dataframe, as shown below.

Table 11: Genre-specific beta values

df_beta[1:10]		df_beta[11:20]	
genre	beta_k	genre	beta_k
Comedy	-0.0022244	War	0.0017305
Romance	-0.0035617	Animation	-0.0153201
Action	-0.0126152	Musical	-0.0102387
Crime	0.0079932	Western	-0.0066510
Thriller	-0.0047380	Mystery	0.0138908
Drama	0.0108468	Film-Noir	0.0306395
Sci-Fi	-0.0119931	Horror	0.0066441
Adventure	-0.0148381	Documentary	0.0624559
Children	-0.0241908	IMAX	-0.0016830
Fantasy	-0.0055682	(no genres listed)	0.1164495

The summation within this model's prediction equation is represented below:

$$\sum_{k=1}^K x_{u,i} \beta_k = x_{u,i}^{\{1\}} \beta_1 + x_{u,i}^{\{2\}} \beta_2 + \dots + x_{u,i}^{\{20\}} \beta_{20}$$

We apply these sums by performing a matrix multiplication between the test set genre columns (`test[,4:23]`) and all 20 β_k values calculated. This multiplication, between a 1799965 per 20 matrix and a 20 per 1 vector, results in a 1799965 per 1 vector. We then add this as a column to the test set, so we can calculate the predictions.

```
sum_x_beta <- as.matrix(test[,4:23])%*%as.matrix(df_beta)
sum_x_beta <- data.frame(sum_x_beta=sum_x_beta[,1])

test <- data.frame(test, sum_x_beta)
rm(df_beta,beta_k,sum_x_beta,all_genres)

predicted_ratings <- test %>%
  mutate(pred = mu + b_i + b_u + sum_x_beta) %>%
  pull(pred)

summary(predicted_ratings)
```

```
>>   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
>> -0.7183  3.1260  3.5581  3.5058  3.9406  6.0652
```

Analyzing the values obtained for ratings, we see there are 216 values under 0.5 and 4157 predicted ratings over 5.

Table 12: RMSE results with the Movie, User and Genre effects model

Method	RMSE
Just the Average	1.0600131
Movie effect model	0.9436204
Movie + User effects model	0.8662966
Movie + User + Genre effects model	0.8661932

While the work needed to obtain this prediction model was arduous, there was very little improvement from the last model.

2.5.5 Regularized Movie + User effects model

When a movie receives ratings from just a few users or when users give only very few reviews, we have more uncertainty. Therefore, larger estimates of both b_i and b_u , negative or positive, are more likely. Consider a case in which we have a movie with 100 user ratings and another movie with just one. While the first movie's b_i will be pretty accurate, the estimate for the second movie will simply be the observed deviation from the average rating, which is a clear sign of overtraining. In this cases, it is better to make this prediction with just the overall average rating μ . That is why we need a form of penalization. Regularization permits us to penalize large estimates that are formed using these small sample sizes.

This time, we will add regularization to one of the previously tried models. Since adding genre didn't have much effect on RMSE and both optimal λ determination (which we'll perform in this section) and β_k calculations are computationally costly, we will use the Movie + User effects model. So now, instead of minimizing the least squares equation, we minimize an equation that adds a penalty. We need to apply regularization to both estimate user and movie effects. We are minimizing:

$$\sum_{u,i} (y_{u,i} - \mu - b_i - b_u)^2 + \lambda \left(\sum_i b_i^2 + \sum_u b_u^2 \right)$$

The formula below shows the new form of calculating b_i , with the inclusion of λ in the mean values calculation. When our sample size n_i is very large, a case which will give us a stable estimate, then the penalty λ is effectively ignored since $n_i + \lambda \approx n_i$. However, when the n_i is small, then the estimate $\hat{b}_i(\lambda)$ is shrunken towards 0. The larger λ , the more we shrink.

$$\hat{b}_i(\lambda) = \frac{1}{\lambda + n_i} \sum_{u=1}^{n_i} (Y_{u,i} - \hat{\mu})$$

In the same way, $\hat{b}_u(\lambda)$ is the regularized b_i by including λ in the division denominator.

$$\hat{b}_u(\lambda) = \frac{1}{\lambda + n_i} \sum_{u=1}^{n_i} (Y_{u,i} - \hat{b}_i - \hat{\mu})$$

We must first calculate various predictions and their respective RMSEs by testing λ values varying from 0 to 10, with a 0.25 stepsize, to determine which regularization value gives us the lowest RMSE.

```
train <- train %>% select(userId, movieId, rating)
test <- test %>% select(userId, movieId, rating)

lambdas <- seq(0, 10, 0.25)
rmsees <- sapply(lambdas, function(l){
  b_i <- train %>%
    group_by(movieId) %>%
    summarize(b_i = sum(rating - mu)/(n()+1))
  b_u <- train %>%
    left_join(b_i, by="movieId") %>%
    group_by(userId) %>%
    summarize(b_u = sum(rating - b_i - mu)/(n()+1))
  predicted_ratings <-
```

```

test %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  mutate(pred = mu + b_i + b_u) %>%
  pull(pred)
return(RMSE(test$rating, predicted_ratings))
})
qplot(lambdas, rmses, xlab='lambda', ylab='RMSE')

```




Figure 13: RMSE values per lambda

By analyzing the plot and retrieving values from the table generated, we can confirm that the optimal λ value is 5, and its corresponding RMSE is 0.8655751.

```

>>   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
>> -0.5784  3.1407  3.5624  3.5100  3.9346  5.9789

```

As usual, the `summary` function shows us that there are predicted values under 0.5 (137 predictions) and over 5.0 (2671).

Table 13: RMSE results with the Regularized Movie and User effects model

Method	RMSE
Just the Average	1.0600131
Movie effect model	0.9436204
Movie + User effects model	0.8662966
Movie + User + Genre effects model	0.8661932
Regularized Movie + User effects model	0.8655751

Once more there is improvement, but not by much.

2.5.6 Matrix Factorization with *reco*system package

[Reco](#)system introduction page

The approach we were taking was not getting much better, so we try a different one: matrix factorization. It is a class of collaborative filtering algorithms used in recommendation systems. Matrix factorization algorithms work by decomposing the user-item interaction matrix into the product of two lower dimensionality rectangular matrices. This family of methods became widely known during the Netflix prize challenge due to its effectiveness. Prediction results can be improved by assigning different regularization weights to the latent factors, based on movies' popularity and users' activeness.

Groups of movies, as well as groups of users, have similar rating patterns. By looking at the correlation between movies, we can see rating patterns, and this is basically what matrix factorization does. It is typically used to approximate an incomplete matrix using the product of two matrices in a latent space. This incomplete matrix, $R_{m \times n}$, is a matrix where rows represent different users and columns represent movies, each cell inside the matrix representing a rating (or the absence of one).

The idea of matrix factorization is to approximate the whole rating matrix $R_{m \times n}$ by the product of two matrices of lower dimensions, $P_{k \times m}$ and $Q_{k \times n}$, such that

$$R \approx P'Q$$

Let p_u be the u -th column of P , and q_i be the i -th column of Q , then the rating given by user u on item i would be predicted as $p'_u q_i$. A typical solution for P and Q is the following:

$$\min_{P, Q} \sum_{(u, i) \in R} [f(p_u, q_i; r_{u, i}) + \mu_P \|p_u\|_1 + \mu_Q \|q_i\|_1 + \frac{\lambda_P}{2} \|p_u\|_2^2 + \frac{\lambda_Q}{2} \|q_i\|_2^2]$$

where (u, i) are locations of observed entries in R , $r_{u, i}$ is the observed rating, f is the loss function, and $\mu_P, \mu_Q, \lambda_P, \lambda_Q$ are penalty parameters to avoid overfitting.

The *reco*system package easily performs the factorization. But to do it, we need to input a parse matrix in triplet form, i.e., each line in the file contains three numbers: `user_index`, `item_index` and `rating`. User and item indexes may start with either 0 or 1, and this can be specified by the `index1` parameter in `data_memory()`. To train and predict, we need to provide objects of class `DataSource`, such as the previously mentioned `data_memory()`, which saves a dataset as an R objects.

```
train_dm <- data_memory(user_index = train$userId, item_index = train$movieId,
                        rating = train$rating, index1 = TRUE)
rm(train)

test_dm <- data_memory(user_index = test$userId, item_index = test$movieId,
                      index1 = TRUE)
test_rating <- test$rating
rm(test)

set.seed(123, sample.kind = "Rounding")
r <- Reco()
params = r$tune(train_dm, opts = list(dim = c(15, 20),
                                       costp_l1 = 0,
                                       costp_l2 = c(0.01, 0.1),
                                       costq_l1 = 0,
                                       costq_l2 = c(0.01, 0.1),
                                       lrate = c(0.075, 0.1), nthread = 2))
```

```
optimal_params = params$min
```

```
r$train(train_dm, opts = c(optimal_params, nthread = 1, niter = 20))
```

20, 0

with the trained model, it is possible to predict the ratings for the test set.

```
predicted_ratings = r$predict(test_dm, out_memory())
```

```
summary(predicted_ratings)
```

```
>>   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
>> -1.594   3.057   3.543   3.472   3.963   6.300
```

object of Output class out_memory(): Result should be returned as R objects
859 and 6221.

```
>> [1] 0.7951871
```

Table 14: RMSE results with the Matrix Factorization Model

Method	RMSE
Just the Average	1.0600131
Movie effect model	0.9436204
Movie + User effects model	0.8662966
Movie + User + Genre effects model	0.8661932
Regularized Movie + User effects model	0.8655751
Matrix Factorization Model	0.7951871

3 Results

3.1 Training full edx dataset with Matrix Factorization

To guarantee the prediction model can be applied to the whole validation set, we will train with full `edx` set.

```
train_dm <- data_memory(user_index = edx$userId, item_index = edx$movieId,
                        rating = edx$rating, index1 = TRUE)
rm(edx)

r$train(train_dm, opts = c(optimal_params, nthread = 1, niter = 20))
```

3.2 Validation set's predicted ratings and RMSE

predicting for validation dataset, we get the following values:

```
validation_dm <- data_memory(user_index = validation$userId,
                             item_index = validation$movieId, index1 = TRUE)
validation_rating <- validation$rating

predicted_ratings = r$predict(validation_dm, out_memory())
rm(train_dm, validation)

summary(predicted_ratings)
```

```
>>   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
>> -1.883  3.058   3.545   3.473   3.965   6.159
```

433 and 3586.

```
>> [1] 0.7887703
```

Table 15: RMSE result for the validation set with Matrix Factorization Model

Method	Validation RMSE
Matrix Factorization Model	0.7887703

4 Conclusion

We visualized the importance of variability between different movies and users, the genre and genre combinations effect on ratings, some not that significant. A lot of models were evaluated, the genre addition didn't have a great effect on the model so we didn't utilize it during the regularized trial, and the best and easiest model to apply was Matrix Factorization with *recosystem*.